

05. CNNs

In this notebook you will learn how to build Convolutional Neural Networks (CNNs) for image processing.

Imports

```
In [1]: %matplotlib inline
```

```
In [2]: import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import sklearn
import sys
import tensorflow as tf
from tensorflow import keras
import time
```

```
In [ ]: print("python", sys.version)
for module in mpl, np, pd, sklearn, tf, keras:
    print(module.__name__, module.__version__)
```

```
In [3]: assert sys.version_info >= (3, 5) # Python ≥3.5 required
assert tf.__version__ >= "2.0" # TensorFlow ≥2.0 required
```

Exercise 1 – Simple CNN


1.1)

Load CIFAR10 using `keras.datasets.cifar10.load_data()`, and split it into a training set (45,000 images), a validation set (5,000 images) and a test set (10,000 images). Make sure the pixel values range from 0 to 1. Visualize a few images using `plt.imshow()`.

```
In [4]: classes = [
    "airplane",
    "automobile",
    "bird",
    "cat",
    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]
```

```
In [5]: (X_train_full, y_train_full), (X_test, y_test) = keras.datasets.cifar10.load_data()
X_train = X_train_full[:5000] / 255
y_train = y_train_full[:5000]
X_valid = X_train_full[5000:] / 255
y_valid = y_train_full[5000:]
X_test = X_test / 255
```

```
In [6]: plt.figure(figsize=(10, 7))
n_rows, n_cols = 10, 15
for row in range(n_rows):
    for col in range(n_cols):
        i = row * n_cols + col
        plt.subplot(n_rows, n_cols, i + 1)
        plt.axis("off")
        plt.imshow(X_train[i])
```



Let's print the classes of the images in the first row:

```
In [7]: for i in range(n_cols):
        print(classes[y_train[i][0]], end=" ")
```

frog truck truck deer automobile automobile bird horse ship cat deer horse horse bird truck

1.2)

Build and train a baseline model with a few dense layers, and plot the learning curves. Use the model's `summary()` method to count the number of parameters in this model.

Tip:

- Recall that to plot the learning curves, you can simply create a Pandas `DataFrame` with the `history.history` dict, then call its `plot()` method.

```
In [39]: model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[32, 32, 3]),
    keras.layers.Dense(64, activation="selu"),
    keras.layers.Dense(64, activation="selu"),
    keras.layers.Dense(64, activation="selu"),
    keras.layers.Dense(10, activation="softmax")
])
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=keras.optimizers.SGD(lr=0.01), metrics=["accuracy"])
history = model.fit(X_train, y_train, epochs=20,
                  validation_data=(X_valid, y_valid))
```

Train on 45000 samples, validate on 5000 samples


Epoch 1/20

45000/45000=====] - 6s 130us/sample - loss: 1.9075 -
acc: 0.3125 - val_loss: 1.8396 - val_acc: 0.3474

Epoch 20/20

45000/45000=====] - 5s 108us/sample - loss: 1.2987 -
acc: 0.5350 - val_loss: 1.4276 - val_acc: 0.4950

```
In [7]: pd.DataFrame(history.history).plot()
plt.axis([0, 19, 0, 1])
plt.show()
```



```
In [8]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 64)	196672
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 64)	4160
dense_3 (Dense)	(None, 10)	650
Total params: 205,642		
Trainable params: 205,642		
Non-trainable params: 0		

1.3)

Build and train a Convolutional Neural Network using a "classical" architecture: $N * (\text{Conv2D} \rightarrow \text{Conv2D} \rightarrow \text{Pool2D}) \rightarrow \text{Flatten} \rightarrow \text{Dense} \rightarrow \text{Dense}$. Before you print the `summary()`, try to manually calculate the number of parameters in your model's architecture, as well as the shape of the inputs and outputs of each layer. Next, plot the learning curves and compare the performance with the previous model.

```
In [1]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
"relu", input_shape=[32, 32, 3]),
    keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(filters=64, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.Conv2D(filters=64, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=keras.optimizers.SGD(lr=0.01), metrics=["accuracy"])
history = model.fit(X_train, y_train, epochs=20,
                    validation_data=(X_valid, y_valid))
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/20

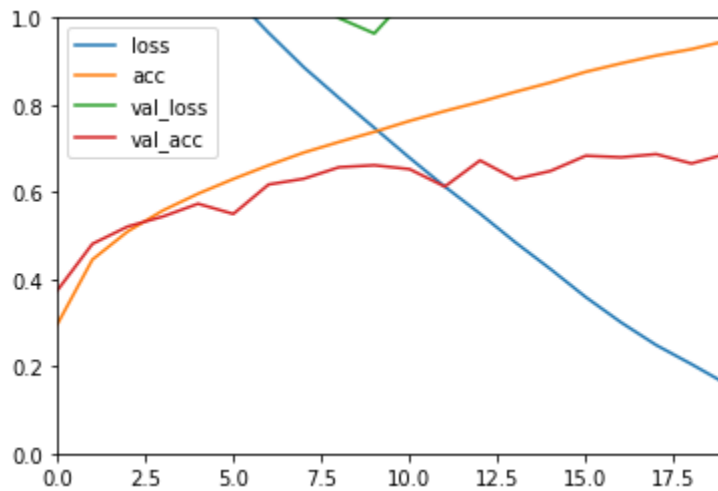
45000/45000=====] - 9s 191us/sample - loss: 1.9518 -
acc: 0.2957 - val_loss: 1.7348 - val_acc: 0.3736

⋮

Epoch 20/20

45000/45000=====] - 6s 125us/sample - loss: 0.1591 -
acc: 0.9471 - val_loss: 1.4888 - val_acc: 0.6878

```
In [10]: pd.DataFrame(history.history).plot()
plt.axis([0, 19, 0, 1])
plt.show()
```



```
In [11]: # Number of params in a convolutional layer =
# (kernel_width * kernel_height * channels_in + 1 for bias) * channels_out
(
    (3 * 3 * 3 + 1) * 32 # in: 32x32x3 out: 32x32x32 Conv2D
+ (3 * 3 * 32 + 1) * 32 # in: 32x32x32 out: 32x32x32 Conv2D
+ 0 # in: 32x32x32 out: 16x16x32 MaxPool2D
+ (3 * 3 * 32 + 1) * 64 # in: 16x16x32 out: 16x16x64 Conv2D
+ (3 * 3 * 64 + 1) * 64 # in: 16x16x64 out: 16x16x64 Conv2D
+ 0 # in: 16x16x64 out: 8x8x64 MaxPool2D
+ 0 # in: 8x8x64 out: 4096 Flatten
+ (4096 + 1) * 128 # in: 4096 out: 128 Dense
+ (128 + 1) * 10 # in: 128 out: 10 Dense
)
```

Out[11]: 591274

Let's check:

```
In [12]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_4 (Dense)	(None, 128)	524416
dense_5 (Dense)	(None, 10)	1290
=====		
Total params: 591,274		
Trainable params: 591,274		
Non-trainable params: 0		
=====		

1.4)

Looking at the learning curves, you can see that the model is overfitting. Add a Batch Normalization layer after each convolutional layer. Compare the model's performance and learning curves with the previous model.

Tip: there is no need for an activation function just before the pooling layers.

```
In [2]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
"relu", input_shape=[32, 32, 3]),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(filters=64, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=64, kernel_size=3, padding="same", activation=
"relu"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=keras.optimizers.SGD(lr=0.01), metrics=["accuracy"])
history = model.fit(X_train, y_train, epochs=20,
                  validation_data=(X_valid, y_valid))
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/20

45000/45000=====] - 8s 170us/sample - loss: 1.3353 -

acc: 0.5280 - val_loss: 1.9303 - val_acc: 0.3884

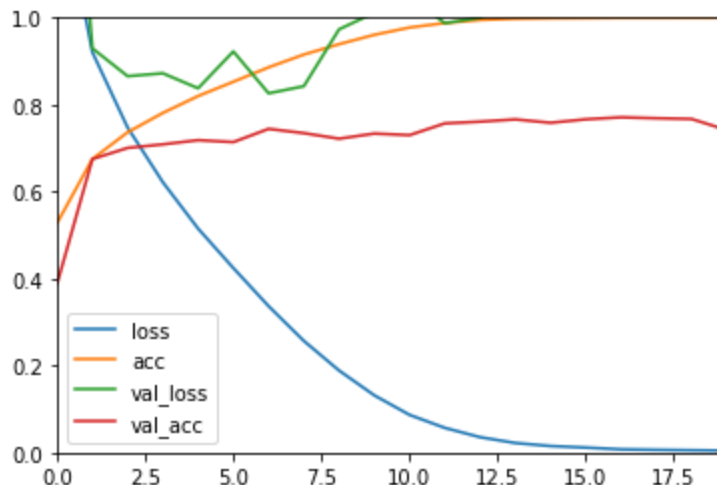
⋮

Epoch 20/20

45000/45000=====] - 7s 158us/sample - loss: 0.0045 -

acc: 1.0000 - val_loss: 1.3410 - val_acc: 0.7418

```
In [14]: pd.DataFrame(history.history).plot()
plt.axis([0, 19, 0, 1])
plt.show()
```



Exercise 2 – Separable Convolutions

2.1)

Replace the `Conv2D` layers with `SeparableConv2D` layers (except the first one), fit your model and compare its performance and learning curves with the previous model.

In [42]:

```
In [3]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
"relu", input_shape=[32, 32, 3]),
    keras.layers.BatchNormalization(),
    keras.layers.SeparableConv2D(filters=32, kernel_size=3, padding="same", ac
tivation="relu"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.SeparableConv2D(filters=64, kernel_size=3, padding="same", ac
tivation="relu"),
    keras.layers.BatchNormalization(),
    keras.layers.SeparableConv2D(filters=64, kernel_size=3, padding="same", ac
tivation="relu"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
model.compile(loss="sparse_categorical_crossentropy",
              optimizer=keras.optimizers.SGD(lr=0.01), metrics=["accuracy"])
history = model.fit(X_train, y_train, epochs=20,
                  validation_data=(X_valid, y_valid))
```

Train on 45000 samples, validate on 5000 samples

Epoch 1/20


45000/45000=====] - 8s 179us/sample - loss: 1.3953 -
acc: 0.5008 - val_loss: 1.7258 - val_acc: 0.4102

:

Epoch 20/20

45000/45000=====] - 7s 163us/sample - loss: 0.0114 -
acc: 0.9995 - val_loss: 1.2084 - val_acc: 0.7588


```
In [16]: pd.DataFrame(history.history).plot()
plt.axis([0, 19, 0, 1])
plt.show()
```



2.2)

Try to estimate the number of parameters in your network, then check your result with `model.summary()`.

Tip: the batch normalization layer adds two parameters for each feature map (the scale and bias).

```
In [19]: # Number of params in a depthwise separable 2D convolution layer =
# kernel_width * kernel_height * channels_in + (channels_in + 1 for bias) * channels_out
(
    (3 * 3 * 3 + 1) * 32          # in: 32x32x3   out: 32x32x32   Conv2D
    + 32 * 2                     # in: 32x32x32  out: 32x32x32   BN
    + 3 * 3 * 32 + (32 + 1) * 32  # in: 32x32x32  out: 32x32x32   SeparableConv2D
    + 32 * 2                     # in: 32x32x32  out: 32x32x32   BN
    + 0                           # in: 32x32x32  out: 16x16x32   MaxPool2D
    + 3 * 3 * 32 + (32 + 1) * 64  # in: 16x16x32  out: 16x16x64   SeparableConv2D
    + 64 * 2                     # in: 16x16x64  out: 16x16x64   BN
    + 3 * 3 * 64 + (64 + 1) * 64  # in: 16x16x64  out: 16x16x64   SeparableConv2D
    + 64 * 2                     # in: 16x16x64  out: 16x16x64   BN
    + 0                           # in: 16x16x64  out: 8x8x64    MaxPool2D
    + 0                           # in: 8x8x64    out: 4096     Flatten
    + (4096 + 1) * 128            # in: 4096     out: 128      Dense
    + (128 + 1) * 10             # in: 128      out: 10       Dense
)
```

Out[19]: 535466

Let's check:

```
In [20]: model.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_v2_4 (Batch Normalization)	(None, 32, 32, 32)	128
separable_conv2d (Separable Conv2D)	(None, 32, 32, 32)	1344
batch_normalization_v2_5 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 32)	0
separable_conv2d_1 (Separable Conv2D)	(None, 16, 16, 64)	2400
batch_normalization_v2_6 (Batch Normalization)	(None, 16, 16, 64)	256
separable_conv2d_2 (Separable Conv2D)	(None, 16, 16, 64)	4736
batch_normalization_v2_7 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_3 (Flatten)	(None, 4096)	0
dense_8 (Dense)	(None, 128)	524416
dense_9 (Dense)	(None, 10)	1290
Total params: 535,850		
Trainable params: 535,466		
Non-trainable params: 384		

Exercise 3 – Pretrained CNNs

3.1)

Using `keras.preprocessing.image.load_img()` followed by `keras.preprocessing.image.img_to_array()`, load one or more images (e.g., `fig.jpg` or `ostrich.jpg` in the `images` folder). You should set `target_size=(299, 299)` when calling `load_img()`, as this is the shape that the Xception network expects.

```
In [22]: img_fig_path = os.path.join("images", "fig.jpg")
img_fig = keras.preprocessing.image.load_img(img_fig_path, target_size=(299, 299))
img_fig = keras.preprocessing.image.img_to_array(img_fig)
```

```
In [23]: plt.imshow(img_fig / 255)
plt.axis("off")
plt.show()
img_fig.shape
```



```
Out[23]: (299, 299, 3)
```

```
In [24]: img_ostrich_path = os.path.join("images", "ostrich.jpg")
img_ostrich = keras.preprocessing.image.load_img(img_ostrich_path, target_size
=(299, 299))
img_ostrich = keras.preprocessing.image.img_to_array(img_ostrich)
```

```
In [25]: plt.imshow(img_ostrich / 255)
plt.axis("off")
plt.show()
img_ostrich.shape
```



```
Out[25]: (299, 299, 3)
```

3.2)

Create a batch containing the image(s) you just loaded, and preprocess this batch using `keras.applications.xception.preprocess_input()`. Verify that the features now vary from -1 to 1: this is what the Xception network expects.

```
In [26]: X_batch = np.array([img_fig, img_ostrich])
X_preproc = keras.applications.xception.preprocess_input(X_batch)
```

```
In [27]: X_preproc.min(), X_preproc.max()
```

```
Out[27]: (-1.0, 1.0)
```

3.3)

Create an instance of the Xception model (`keras.applications.xception.Xception`) and use its `predict()` method to classify the images in the batch. You can use `keras.applications.resnet50.decode_predictions()` to convert the output matrix into a list of top-N predictions (with their corresponding class labels).

```
In [28]: model = keras.applications.xception.Xception()
Y_proba = model.predict(X_preproc)
```

```
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.4/xception_weights_tf_dim_ordering_tf_kernels.h5
91889664/91884032=====] - 3s 0us/step
```

```
In [29]: Y_proba.shape
```

```
Out[29]: (2, 1000)
```

```
In [30]: np.argmax(Y_proba, axis=1)
```

```
Out[30]: array([952,   9])
```

```
In [31]: decoded_predictions = keras.applications.resnet50.decode_predictions(Y_proba)
for preds in decoded_predictions:
    for wordnet_id, name, proba in preds:
        print("{} ({}): {:.1f}%".format(name, wordnet_id, 100 * proba))
    print()
```

```
Downloading data from https://s3.amazonaws.com/deep-learning-models/image-models/imagenet_class_index.json
40960/35363=====] - 0s 2us/step
```

```
fig (n07753113): 99.9%
grocery_store (n03461385): 0.0%
mushroom (n07734744): 0.0%
butternut_squash (n07717556): 0.0%
pomegranate (n07768694): 0.0%
```

```
ostrich (n01518878): 98.8%
bustard (n02018795): 0.1%
black_swan (n01860187): 0.0%
white_stork (n02002556): 0.0%
cock (n01514668): 0.0%
```

Exercise 4 – Data Augmentation and Transfer Learning

In this exercise you will reuse a pretrained Xception model to build a flower classifier.

First, let's download the dataset:

```
In [32]: import tensorflow as tf
         from tensorflow import keras
         import os

         flowers_url = "http://download.tensorflow.org/example_images/flower_photos.tgz"
         flowers_path = keras.utils.get_file("flowers.tgz", flowers_url, extract=True)
         flowers_dir = os.path.join(os.path.dirname(flowers_path), "flower_photos")

         Downloading data from http://download.tensorflow.org/example_images/flower_photos.tgz
         228818944/228813984=====] - 4s 0us/step
```

```
In [33]: for root, subdirs, files in os.walk(flowers_dir):
         print(root)
         for filename in files[:3]:
             print(" ", filename)
         if len(files) > 3:
             print(" ...")

/home/jupyter/.keras/datasets/flower_photos
  LICENSE.txt
/home/jupyter/.keras/datasets/flower_photos/roses
  3664842094_5fd60ee26b.jpg
  459042023_6273adc312_n.jpg
  5492988531_574cdc2bf0_n.jpg
  ...
/home/jupyter/.keras/datasets/flower_photos/dandelion
  515143813_b3afb08bf9.jpg
  14085038920_2ee4ce8a8d.jpg
  3696596109_4c4419128a_m.jpg
  ...
/home/jupyter/.keras/datasets/flower_photos/sunflowers
  6606823367_e89dc52a95_n.jpg
  5998488415_a6bacd9f83.jpg
  678714585_addc9aaef.jpg
  ...
/home/jupyter/.keras/datasets/flower_photos/tulips
  16169741783_deeab1a679_m.jpg
  13539384593_23449f7332_n.jpg
  16711791713_e54bc9claf_n.jpg
  ...
/home/jupyter/.keras/datasets/flower_photos/daisy
  9345273630_af3550031d.jpg
  14088053307_1a13a0bf91_n.jpg
  4746633946_23933c0810.jpg
  ...
```

4.1)

Build a `keras.preprocessing.image.ImageDataGenerator` that will preprocess the images and do some data augmentation (the [documentation](#) contains useful examples):

- It should at least perform horizontal flips and keep 10% of the data for validation, but you may also make it perform a bit of rotation, rescaling, etc.
- Also make sure to apply the Xception preprocessing function (using the `preprocessing_function` argument).
- Call this generator's `flow_from_directory()` method to get an iterator that will load and preprocess the flower photos from the `flower_photos` directory, setting the target size to (299, 299) and `subset` to "training".
- Call this method again with the same parameters except `subset="validation"` to get a second iterator for validation.
- Get the next batch from the validation iterator and display the first image from the batch.

```
In [34]: datagen = keras.preprocessing.image.ImageDataGenerator(
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        validation_split=0.1,
        preprocessing_function=keras.applications.xception.preprocess_input)


train_generator = datagen.flow_from_directory(
    flowers_dir,
    target_size=(299, 299),
    batch_size=32,
    subset="training")

valid_generator = datagen.flow_from_directory(
    flowers_dir,
    target_size=(299, 299),
    batch_size=32,
    subset="validation")
```

Found 3306 images belonging to 5 classes.

Found 364 images belonging to 5 classes.

```
In [35]: X_batch, y_batch = next(valid_generator)
plt.imshow((X_batch[0] + 1)/2)
plt.axis("off")
plt.show()
```



4.2)

Now let's build the model:

- Create a new `Xception` model, but this time set `include_top=False` to get the model without the top layer. **Tip:** you will need to access its `input` and `output` properties.
- Make all its layers non-trainable.
- Using the functional API, add a `GlobalAveragePooling2D` layer (feeding it the Xception model's output), and add a `Dense` layer with 5 neurons and the Softmax activation function.
- Compile the model. **Tip:** don't forget to add the "accuracy" metric.
- Fit your model using `fit_generator()`, passing it the training and validation iterators (and setting `steps_per_epoch` and `validation_steps` appropriately).

```
In [36]: n_classes = 5

base_model = keras.applications.xception.Xception(include_top=False)

for layer in base_model.layers:
    layer.trainable = False

global_pool = keras.layers.GlobalAveragePooling2D()(base_model.output)
predictions = keras.layers.Dense(n_classes, activation='softmax')(global_pool)

model = keras.models.Model(base_model.input, predictions)

model.compile(loss="categorical_crossentropy",
              optimizer="sgd", metrics=["accuracy"])
```

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.4/xception_weights_tf_dim_ordering_tf_kernels_notop.h5
83689472/83683744=====] - 3s 0us/step

```
In [38]: history = model.fit_generator(
    train_generator,
    steps_per_epoch=3306 // 32,
    epochs=50,
    validation_data=valid_generator,
    validation_steps=364 // 32)
```

```
In [ ]: pd.DataFrame(history.history).plot()
plt.axis([0, 19, 0, 1])
plt.show()
```

Object Detection Project

The Google [Street View House Numbers](#) (SVHN) dataset contains pictures of digits in all shapes and colors, taken by the Google Street View cars. The goal is to classify and locate all the digits in large images.

- Train a Fully Convolutional Network on the 32x32 images.
- Use this FCN to build a digit detector in the large images.